


# MSAI-437: Deep Learning

## Homework 2 (Group 9)

### Convolutional Neural Network


#### Architecture




#### Hyperparameters

- **Loss Function:** Binary Cross-Entropy
- **Optimizer:** Adam Optimizer
- **Learning Rate:** 0.001
- **Number of Training Epochs:** 20

## Learning Curves



## Validation Set Performance



**Validation accuracy = 0.9549**

## Limitations

The main limitation of our CNN is that it is unable to represent spatial information about the features it is extracting. This can cause issues because if certain filters are mainly contributing to the unhealthy classification and that feature is found in the background of a healthy leaf, the image will be classified incorrectly. The CNN has no way of saying “This feature should be found near the center of the image (where the leaf is).”

# Autoencoder

## Methodology


We trained the autoencoder on the Beans dataset using only the healthy leaves. We used the input images themselves as labels, and minimized the Mean Squared Error. The encoder condenses the  $3 \times 500 \times 500$  image into a  $1 \times 1600$  latent space.

Once the autoencoder was trained, we discarded the decoder half of the architecture (see next two pages), retaining only the encoder.


We then used a feed forward neural network (FFNN) to classify the Beans dataset. We passed all the images through the encoder to create a  $1 \times 1600$  vector. It was then just a simple matter of running a FFNN on the  $1 \times 1600$  vector for classification into healthy or unhealthy. We used Binary Cross Entropy as the loss function.

## Architecture


### Encoder



### Feed Forward Neural Network



## Decoder



## Hyperparameters

### Autoencoder


- **Loss Function:** Mean Squared Error
- **Optimizer:** Adam Optimizer
- **Learning Rate:** 0.00001
- **Number of Training Epochs:** 500

### Feed Forward Neural Network


- **Loss Function:** Binary Cross-Entropy
- **Optimizer:** Adam Optimizer
- **Learning Rate:** 0.0001
- **Number of Training Epochs:** 100

## Learning Curves

### Autoencoder

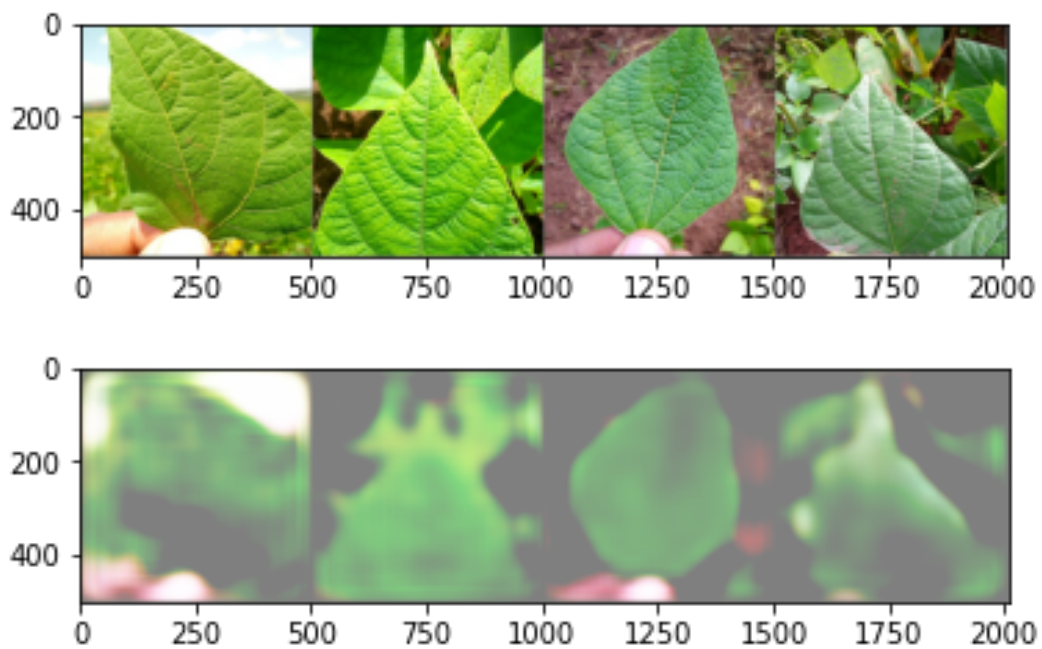


## Feed Forward Neural Network

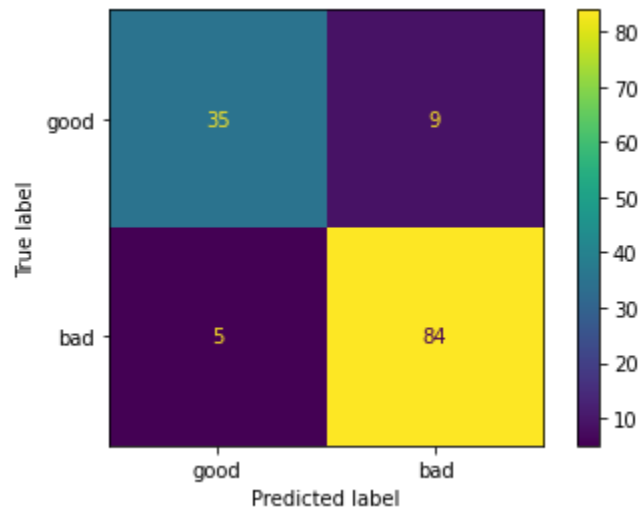


## Validation Set Performance

### Autoencoder Image Reconstruction



## Feed Forward Neural Network Classification



**Validation accuracy = 0.8947**

## Limitations

- The autoencoder was trained on just healthy leaf images so there is no guarantee about a consistent encoding of unhealthy leaves.
- The autoencoder and the feed forward neural network (FFNN) were not being trained simultaneously towards the same task. The FFNN was trying to classify both healthy and unhealthy leaves condensed into the latent space by an autoencoder trained only on healthy leaves. It would have been more efficient to train the autoencoder on both types of leaves.



# Alternate Approaches and Methods

1. **Multiclass classification:** We had two different types of diseases in the dataset, but chose to lump them together. Maybe if we attempt to learn the differences between them, we can perform better overall.
2. **Knowledge-based approach:** We can imagine a knowledge base which knows more about the visual features of each type of disease, and a computer vision algorithm that attempts to extract these features.
3. **Data augmentation:** There are several ways to increase the small amount of training data that we have - we can blur, zoom, rotate, and realign the images to create “new” data points.
4. **Better targeting:** We can first try to locate the leaf of interest in the image (by segmentation / bounding boxes, for example), and then apply classification techniques.
5. **Transfer learning:** We can try to use a classifier pretrained on some other (similar) classification task and fine-tune it for our needs.
6. **K-Nearest Neighbors:** KNN performs well on small data sets, but we will need to convert the images to a latent representation (such as with the autoencoder above) before it would make sense to measure “distance” between images.
7. **Denoising Autoencoder:** We can try to mask portions of the leaf in the input, and train the autoencoder. This way, it learns a more robust representation of the leaves and is more resistant to noise.